

SEC1

**Company
Confidential**

**Access
for Capgemini
users and
authorized
third parties**

Relational Databases

Maciej Lipski

2021-04-06



Agenda



- Before we Begin
- Relational Databases (DB)
- Structured Query Language (SQL)
- Data locking strategies
- Advanced Databases Topics
- Modeling databases
- Assignments

- **Before we Begin**

- Relational Databases (DB)
- Structured Query Language (SQL)
- Data locking strategies
- Advanced Databases Topics
- Modeling databases
- Assignments

Agenda




- Before we Begin
- **Relational Databases (DB)**
- Structured Query Language (SQL)
- Data locking strategies
- Advanced Databases Topics
- Modeling databases
- Assignments

What is this database thing?




- **Database (DB)** – organized collection of data
- **Database Management System (DBMS)** – computer software, allows the definition, creation, querying, update, and administration of databases.
- In our daily work we usually mix those two terms.



Why do we need databases?

Is it really needed? Do we care?





- Abstraction layer between data and application
- It's about speed, scalability, integrity, security, transaction support, availability, consistency... You don't want to implement it yourself.
- NOTE: Our development environments might give an impression, that the database is irrelevant (in-memory database, initialized on applications start and destroyed when application stops). It is not



Who is using DB?

- **Everyone!**
- End-users
- Developers of the applications
- Testers
- Designers of the database
- System analytics
- Business Analytics
- Database Administrators
- ...

Different database types



- **Oracle**
- **MySQL**
- **MS SQL Server**
- PostgreSQL
- IBM Db2
- MariaDB

- InterSystem Caché
- Db4o
- Matisse
- Perst
- ZopeDB

- Cassandra
- HBase
- Accumulo
- Scylla
- HBase

- MongoDB
- Amazon DynamoDB
- Microsoft Azure Cosmos DB
- CouchDB
- DocumentDB
- MemcacheDB

- Redis
- Amazon DynamoDB
- Oracle NoSQL
- Apache Ignite

- Neo4j
- Microsoft Azure Cosmos DB
- OrientDB
- ArangoDB
- AllegroGraph

Source: <https://www.alooma.com/blog/types-of-modern-databases> ; <https://db-engines.com/en/ranking>

Rankings of popularity

URL: <https://db-engines.com/en/ranking>

361 systems in ranking, January 2021

Rank			DBMS	Database Model	Score		
Jan 2021	Dec 2020	Jan 2020			Jan 2021	Dec 2020	Jan 2020
1.	1.	1.	Oracle +	Relational, Multi-model i	1322.93	-2.66	-23.75
2.	2.	2.	MySQL +	Relational, Multi-model i	1252.06	-3.40	-22.60
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model i	1031.23	-6.85	-67.31
4.	4.	4.	PostgreSQL +	Relational, Multi-model i	552.23	+4.65	+45.03
5.	5.	5.	MongoDB +	Document, Multi-model i	457.22	-0.51	+30.26
6.	6.	6.	IBM Db2 +	Relational, Multi-model i	157.17	-3.26	-11.53
7.	7.	↑ 8.	Redis +	Key-value, Multi-model i	155.01	+1.38	+6.26
8.	8.	↓ 7.	Elasticsearch +	Search engine, Multi-model i	151.25	-1.24	-0.19
9.	9.	↑ 10.	SQLite +	Relational	121.89	+0.21	-0.25
10.	10.	↑ 11.	Cassandra +	Wide column	118.08	-0.76	-2.59
11.	11.	↓ 9.	Microsoft Access	Relational	115.33	-1.41	-13.24
12.	12.	↑ 13.	MariaDB +	Relational, Multi-model i	93.79	+0.18	+6.34
13.	13.	↓ 12.	Splunk	Search engine	87.66	+0.66	-1.01
14.	14.	↑ 15.	Teradata +	Relational, Multi-model i	72.59	-1.24	-5.70
15.	↑ 16.	↑ 25.	Microsoft Azure SQL Database	Relational, Multi-model i	71.36	+1.87	+43.16
16.	↓ 15.	↓ 14.	Hive	Relational	70.43	+0.16	-13.81
17.	17.	↓ 16.	Amazon DynamoDB +	Multi-model i	69.14	+0.01	+7.11

Cloud Computing, Cloud Databases



Database as a Service

SaaS

IaaS

PaaS

Relational databases

- Based on the mathematical model of relation (you know, a subset of a cartesian product).
- A database consists of multiple tables.

Table name

employee

Column

Each column has name and a defined type (like string, number, enum, date...)

id	first_name	last_name	hired	email	created_by	created_at
1	Maciej	Kowalski	2017-01-01	mkowalski@capg.com	hr1	2019-02-17
2	Armando	Roggio	2014-03-13	armando@some.com	hr1	2019-01-01
3	Barack	Obama	2017-01-20	barack@usa.com	admin	2019-01-02
4	Bono	Vox	2014-07-13	bono@u2.com	admin	2019-01-01
5	Bill	Gates	2009-02-20	bill@microsoft.com	admin	2019-01-01
6	Steve	Jobs	1988-02-08	steve@apple.com	null	null
7	Larry	Ellison	2000-08-09	larry@oracle.com	admin	2019-01-01
8	Maciej	Nowak	2017-02-01	mnowak@capg.com	hr2	2017-02-01

Row, tuple

Primary Key

- Attribute or set of attributes, which identifies row
- Can be natural or technical (some ,id' field)

first_name	last_name	hired	email	created_by	created_at
Maciej	Kowalski	2017-01-01	mkowalski@capg.com	hr1	2019-02-17
Armando	Roggio	2014-03-13	armando@some.com	hr1	2019-01-01
Barack	Obama	2017-01-20	barack@usa.com	admin	2019-01-02
Bono	Vox	2014-07-13	bono@u2.com	admin	2019-01-01
Bill	Gates	2009-02-20	bill@microsoft.com	admin	2019-01-01
Steve	Jobs	1988-02-08	steve@apple.com	null	null
Larry	Ellison	2000-08-09	larry@oracle.com	admin	2019-01-01
Maciej	Nowak	2017-02-01	mnowak@capg.com	hr2	2017-02-01

Foreign Key

- Attribute or set of attributes, which points to row in other (or same) table

Employee

id	first_name	last_name	position_id	manager_id
1	Maciej	Kowalski	1	null
2	Armando	Roggio	2	1
3	Barack	Obama	2	1
4	Bono	Vox	3	2
5	Bill	Gates	3	2
6	Steve	Jobs	3	2
7	Larry	Ellison	3	3
8	Maciej	Nowak	3	3

Position

id	position
1	CEO
2	Head of department
3	Developer

Mails of employees

employee_id	email
1	kowalski@capg.com
2	roggio@some.com
3	barack@usa.com
1	m.kowalski@test.pl

Normalizations and relations

- Re-structure tables and relations in order to satisfy so called normal forms
- Goal: reduce data duplications, improve integrity etc

id	name	email
1	Maciej Kowalski	mkowalski@capg.com
2	Armando Roggio	armando@some.com
3	Barack Obama	barack@usa.com
4	Bono Vox	bono@u2.com

id	first_name	last_name	email
1	Maciej	Kowalski	mkowalski@capg.com
2	Armando	Roggio	armando@some.com
3	Barack	Obama	barack@usa.com
4	Bono	Vox	bono@u2.com

id	first_name	last_name	address
1	Maciej	Kowalski	Legnicka 48H, 54-202 Wrocław
2	Armando	Roggio	Tęczowa 301, 53-601 Wrocław

id	first_name	last_name	zip_code	city	street	apartment
1	Maciej	Kowalski	54-202	Wrocław	Legnicka	48H
2	Armando	Roggio	53-601	Wrocław	Tęczowa	301

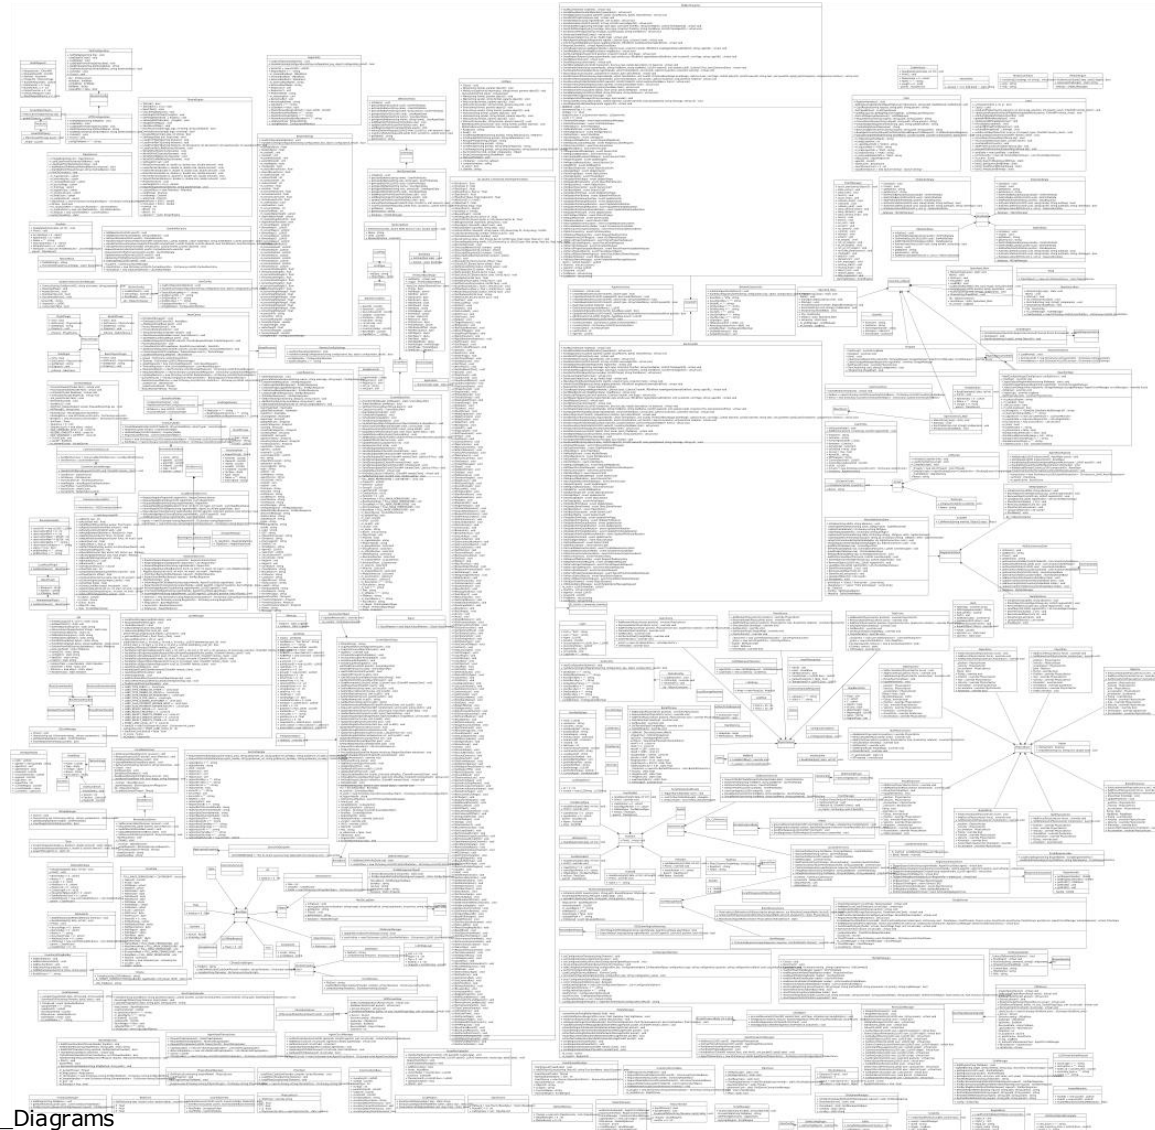
id	last_name	email1	email2
1	Kowalski	kowal@capg.com	mkowalski@test.pl
2	Roggio	roggio@some.com	
3	Obama	barack@usa.com	
4	Vox		

id	last_name
1	Kowalski
2	Roggio
3	Obama
4	Vox

emp_id	email
1	kowal@capg.com
2	roggio@some.com
3	barack@usa.com
1	mkowalski@test.pl

https://en.wikipedia.org/wiki/Database_normalization

Relations... everywhere... relations



Source:http://opensimulator.org/wiki/OpenSim:_Class_Diagrams


Agenda



- Before we Begin
- Relational Databases (DB)
- **Structured Query Language (SQL)**
- Data locking strategies
- Advanced Databases Topics
- Modeling databases
- Assignments

SQL

“Structured
Query
Language”



“S-Q-L”



“Se-quel”




“Skewl”
“Squeal”
“Squiggle”




Source: <https://learnsql.com/blog/sql-or-sequel/>

Structured Query Language (SQL)




Standardized

Current version: [SQL:2016](#)



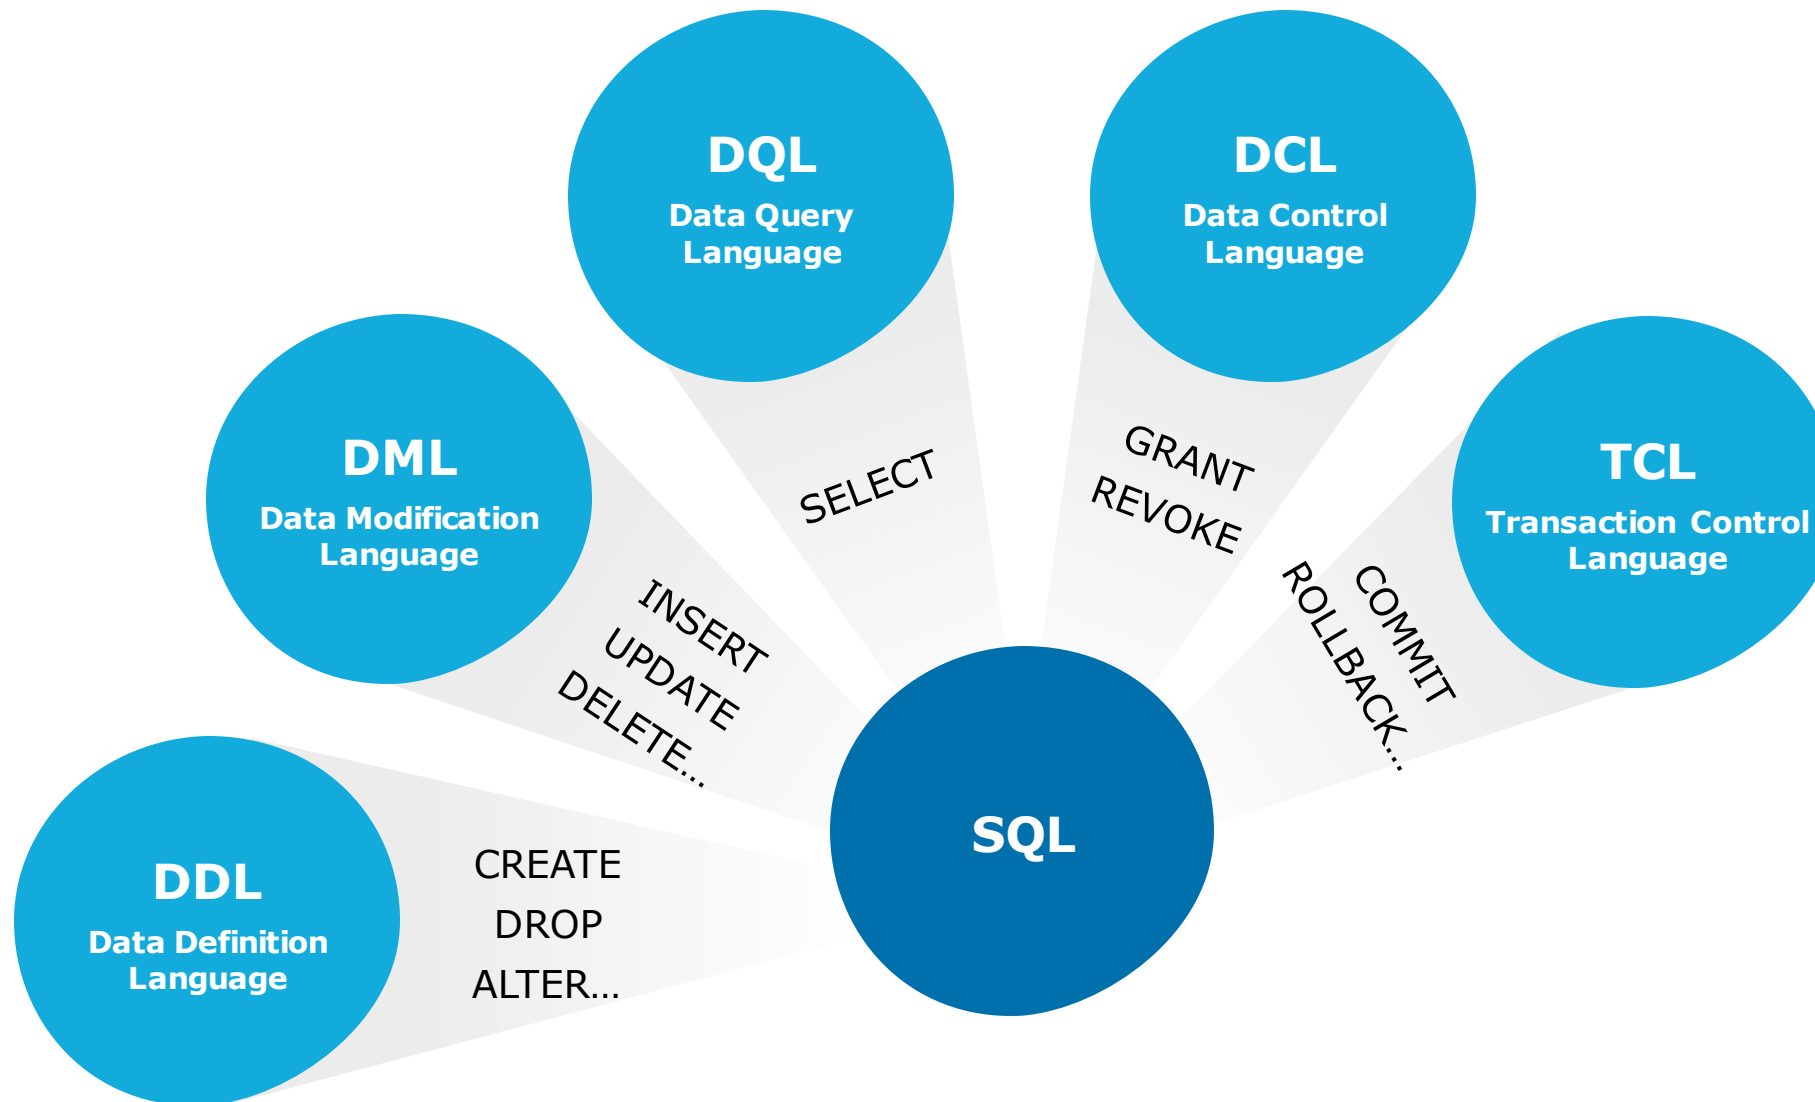
Every vendor adds his own flavors



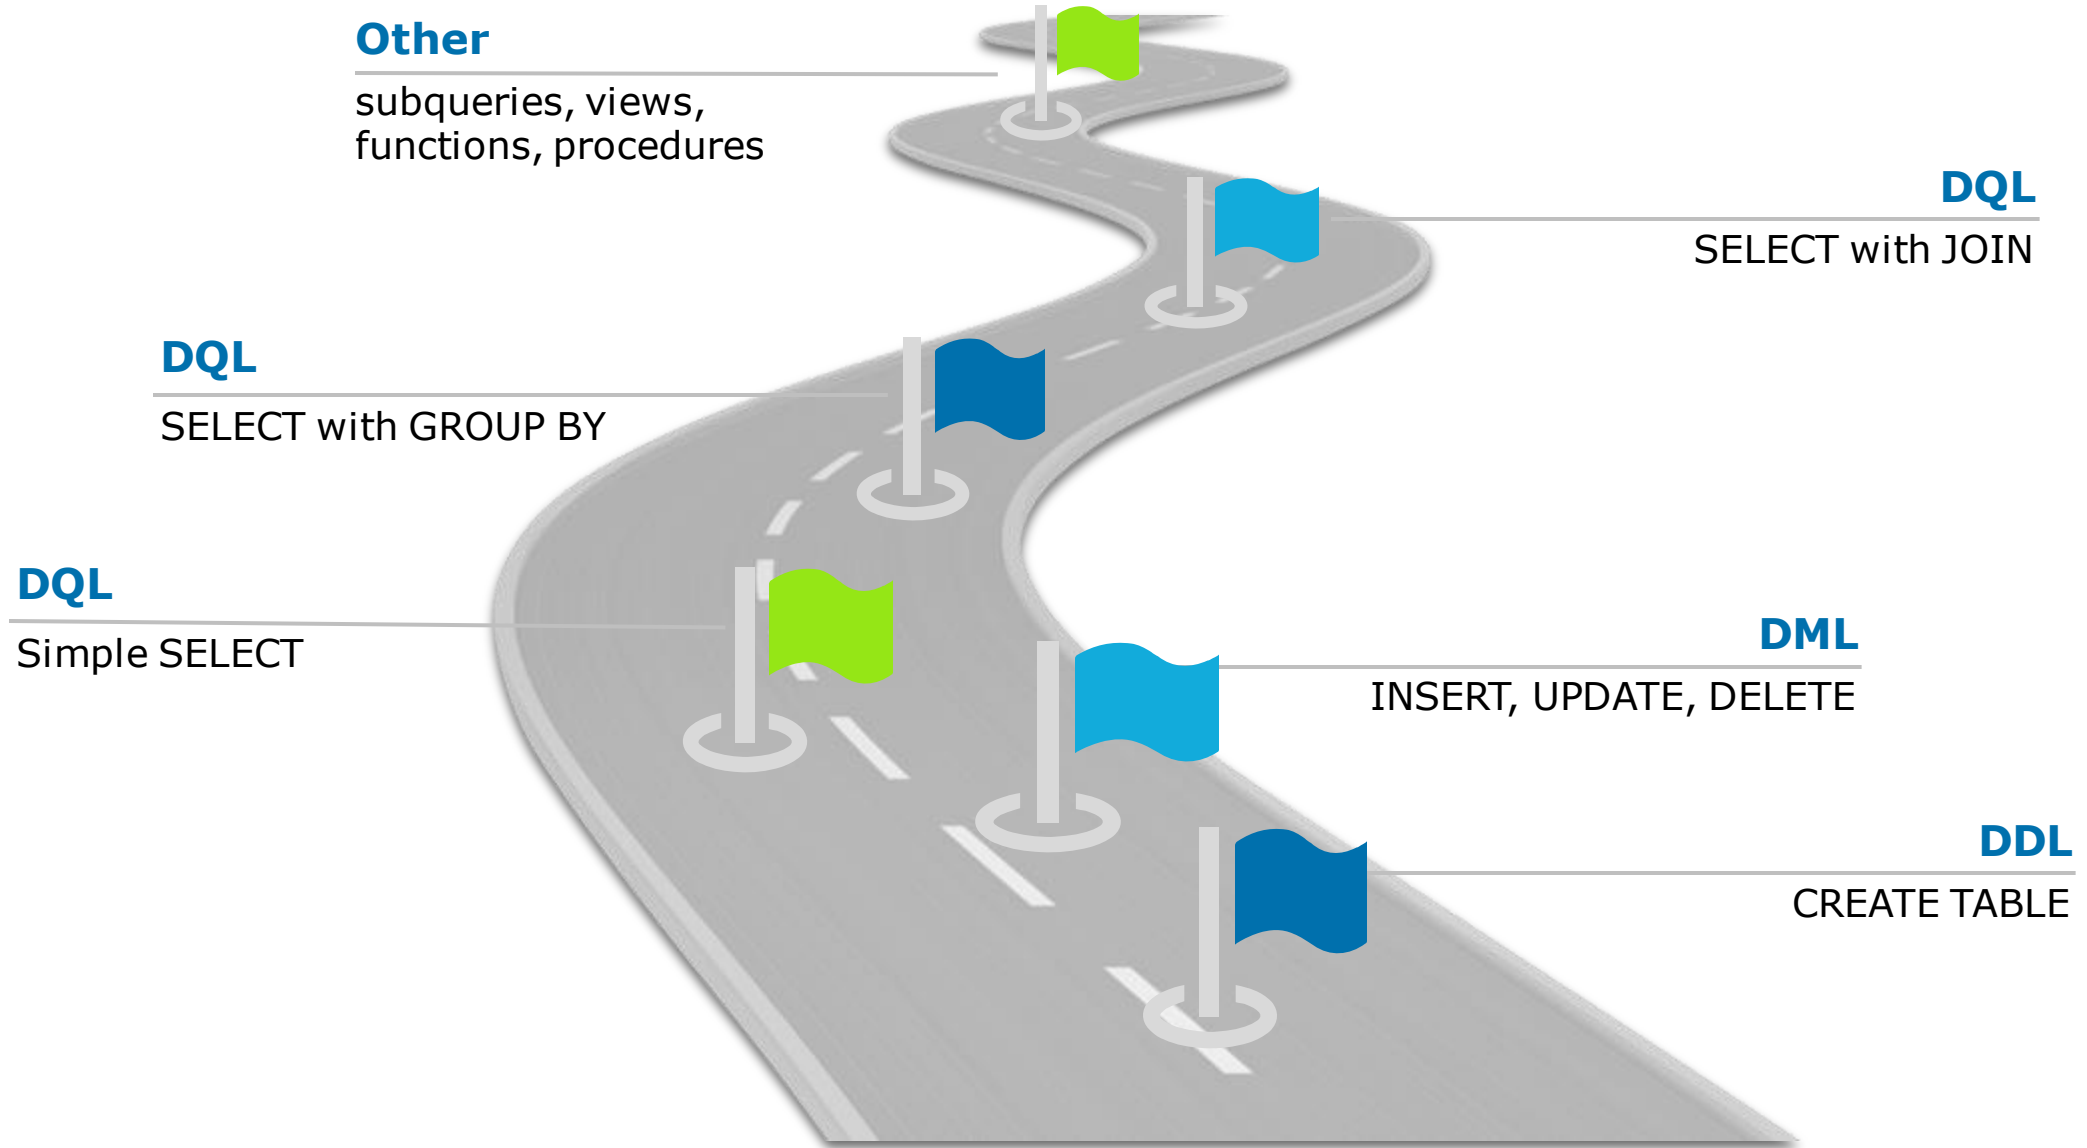
Declarative language

Define how the output should look like.
in contrast with imperative programming,
where you define an algorithm.

SQL commands are categorized into categories



Our trip...



DDL: CREATE TABLE

Table name

```
CREATE TABLE employee (  
    id INT NOT NULL AUTO_INCREMENT,  
    first_name VARCHAR(45) NOT NULL,  
    last_name VARCHAR(45) NOT NULL,  
    hired DATE NOT NULL,  
    mail VARCHAR(50) NOT NULL,  
    created_by VARCHAR(50) NULL,  
    created_at DATETIME NULL  
        DEFAULT CURRENT_TIMESTAMP,  
    PRIMARY KEY (id),  
    UNIQUE INDEX uq_employee_mail (mail ASC)  
);
```

Table column definition:
<name> <type> <options>

Optional constraints

Common types:

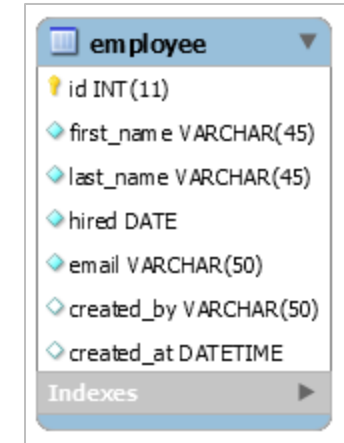
- CHAR
- VARCHAR
- INT
- ENUM
- DATE, DATETIME
- TIMESTAMP...

Common options:

- NULL / NOT NULL
- DEFAULT <value>
- AUTO_INCREMENT
- ON UPDATE...

Significance of Primary Key (PK):

- PK identifies one table row
- PK can span over multiple columns
- Usually, all tables have a PK
- DB usually creates an index for a PK

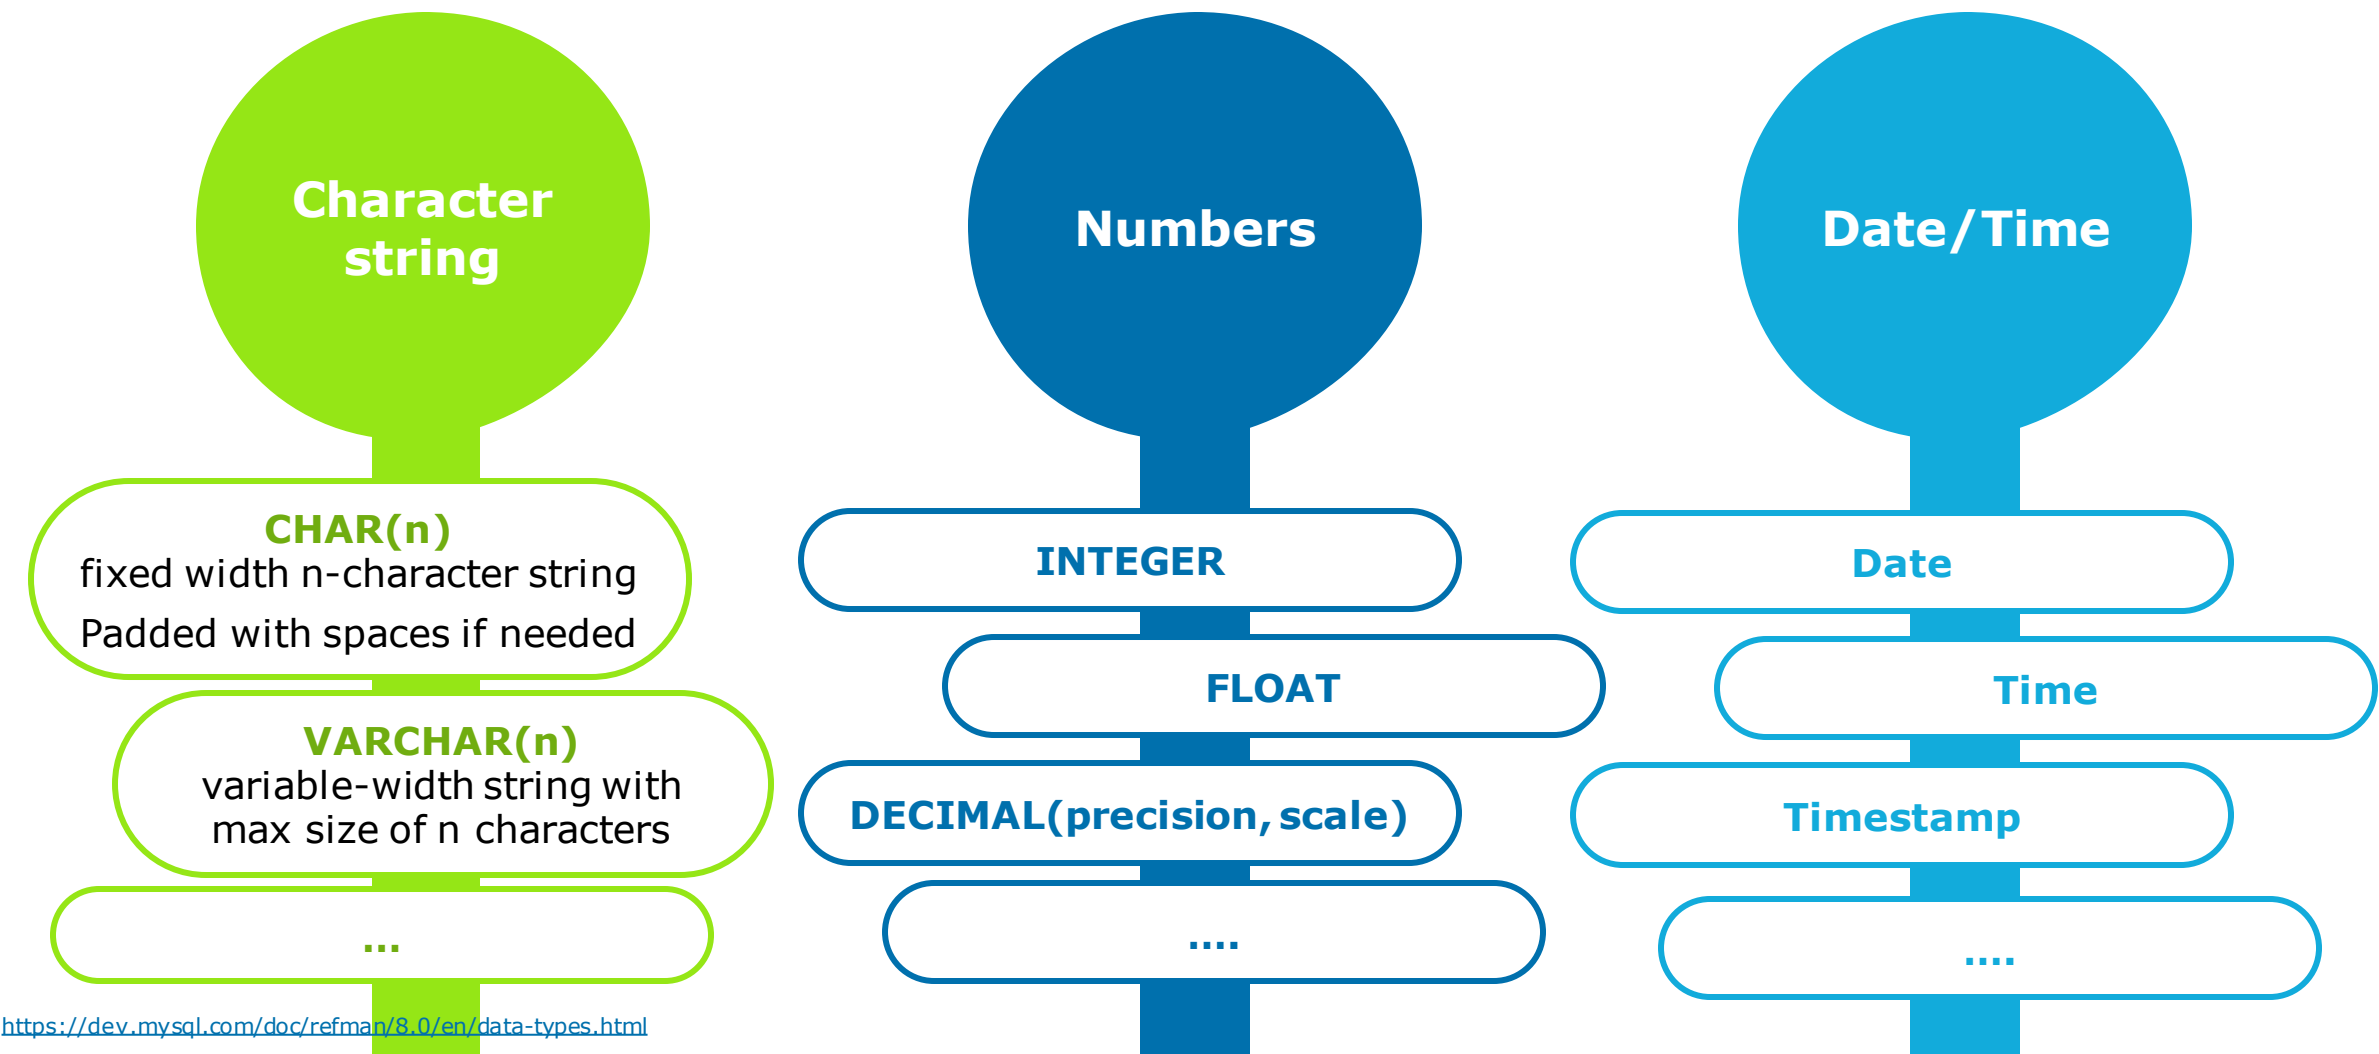


employee	
id	INT(11)
first_name	VARCHAR(45)
last_name	VARCHAR(45)
hired	DATE
email	VARCHAR(50)
created_by	VARCHAR(50)
created_at	DATETIME
Indexes	

<https://dev.mysql.com/doc/refman/8.0/en/sql-data-definition-statements.html>

Data types of columns

Every database can have its own data types



<https://dev.mysql.com/doc/refman/8.0/en/data-types.html>

DML: INSERT INTO

Schema name

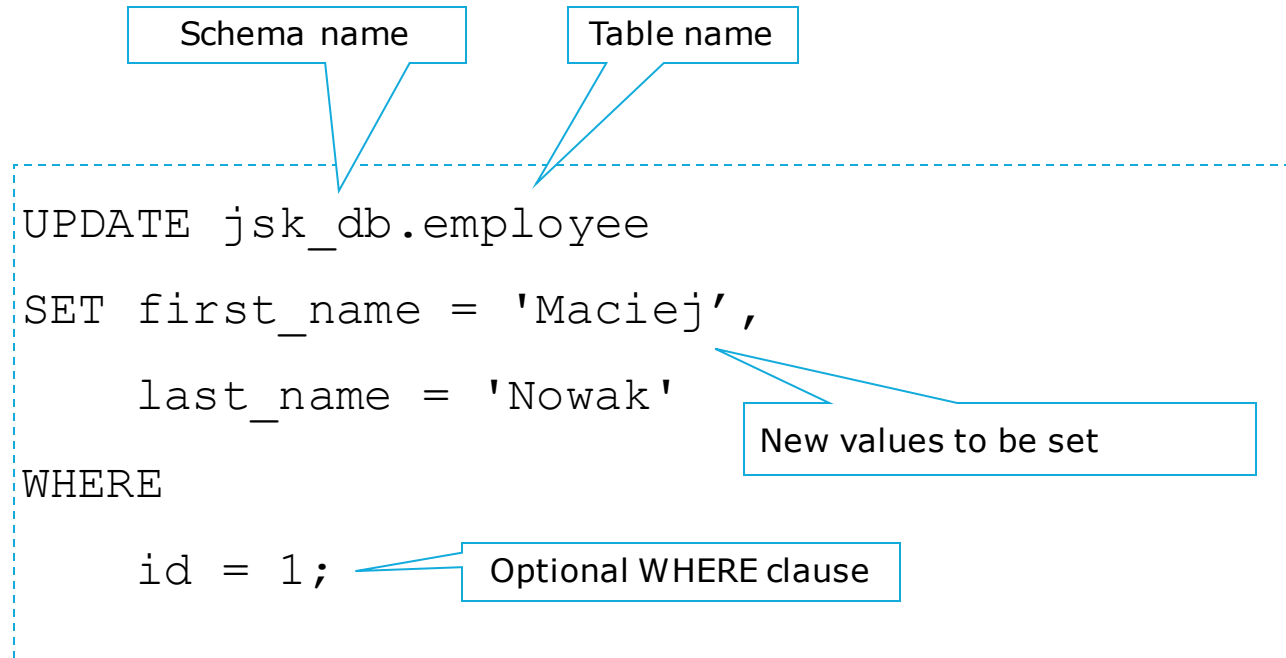
Table name

Optional, comma separated list of columns.
All NOT NULL columns without DEFAULT or AUTO_INCREMENT value must be present.

```
INSERT INTO jsk_db.employee(first_name, last_name, hired, email)
VALUES
    ('Armando', 'Roggio', '2014-03-13', 'armando@some.com'),
    ('Barack', 'Obama', '2017-01-20', 'barack@usa.com'),
    ('Bono', 'Vox', '2014-07-13', 'bono@u2.com', 'admin');
```

Comma separated list of values for each row.

DML: UPDATE



The diagram shows an SQL UPDATE statement enclosed in a dashed blue box. Callout boxes point to specific parts of the statement: 'Schema name' points to 'jsk_db', 'Table name' points to 'employee', 'New values to be set' points to the SET clause, and 'Optional WHERE clause' points to the WHERE clause.

```
UPDATE jsk_db.employee
SET first_name = 'Maciej',
    last_name = 'Nowak'
WHERE
    id = 1;
```

- **DANGER:** Since WHERE clause is optional, running the statement without it will change all rows in the table.
- NOTE: there is no UNDO button in the DB. And it's surprisingly easy to screw things up (BUT there are Transactions)
- When doing difficult migrations it might be wise to store table data as backup (CREATE TABLE ... AS SELECT ...).

DML: DELETE FROM

Schema name

Table name

```
DELETE FROM jsk_db.employee
```

WHERE

```
id = 1;
```

Optional WHERE clause

Let's do some excersises...


Sample table „employee”

id	first_name	last_name	hired	email	created_by	created_at
1	Maciej	Kowalski	2017-01-01	mkowalski@capg.com	hr1	2019-02-17
2	Armando	Roggio	2014-03-13	armando@some.com	hr1	2019-01-01
3	Barack	Obama	2017-01-20	barack@usa.com	admin	2019-01-02
4	Bono	Vox	2014-07-13	bono@u2.com	admin	2019-01-01
5	Bill	Gates	2009-02-20	bill@microsoft.com	admin	2019-01-01
6	Steve	Jobs	1988-02-08	steve@apple.com	null	null
7	Larry	Ellison	2000-08-09	larry@oracle.com	admin	2019-01-01
8	Maciej	Nowak	2017-02-01	mnowak@capg.com	hr2	2017-02-01

<https://dev.mysql.com/doc/refman/8.0/en/select.html>

Understanding GROUP BY

id	first_name	last_name	hired	created_by	created_at
1	Maciej	Kowalski	2017-01-01	hr1	2019-02-17
2	Armando	Roggio	2014-03-13	hr1	2019-01-01
3	Barack	Obama	2017-01-20	admin	2019-01-02
4	Bono	Vox	2014-07-13	admin	2019-01-01
5	Bill	Gates	2009-02-20	admin	2019-01-01
6	Steve	Jobs	1988-02-08	null	null
7	Larry	Ellison	2000-08-09	admin	2019-01-01
8	Maciej	Nowak	2017-02-01	hr2	2017-02-01



```
SELECT e.created_by,
       COUNT(*),
       MAX(e.id)
FROM employee AS e
GROUP BY e.created_by;
```

created_by	COUNT(*)	MAX(e.id)
hr1	2	2
admin	4	7
null	1	6
hr2	1	8

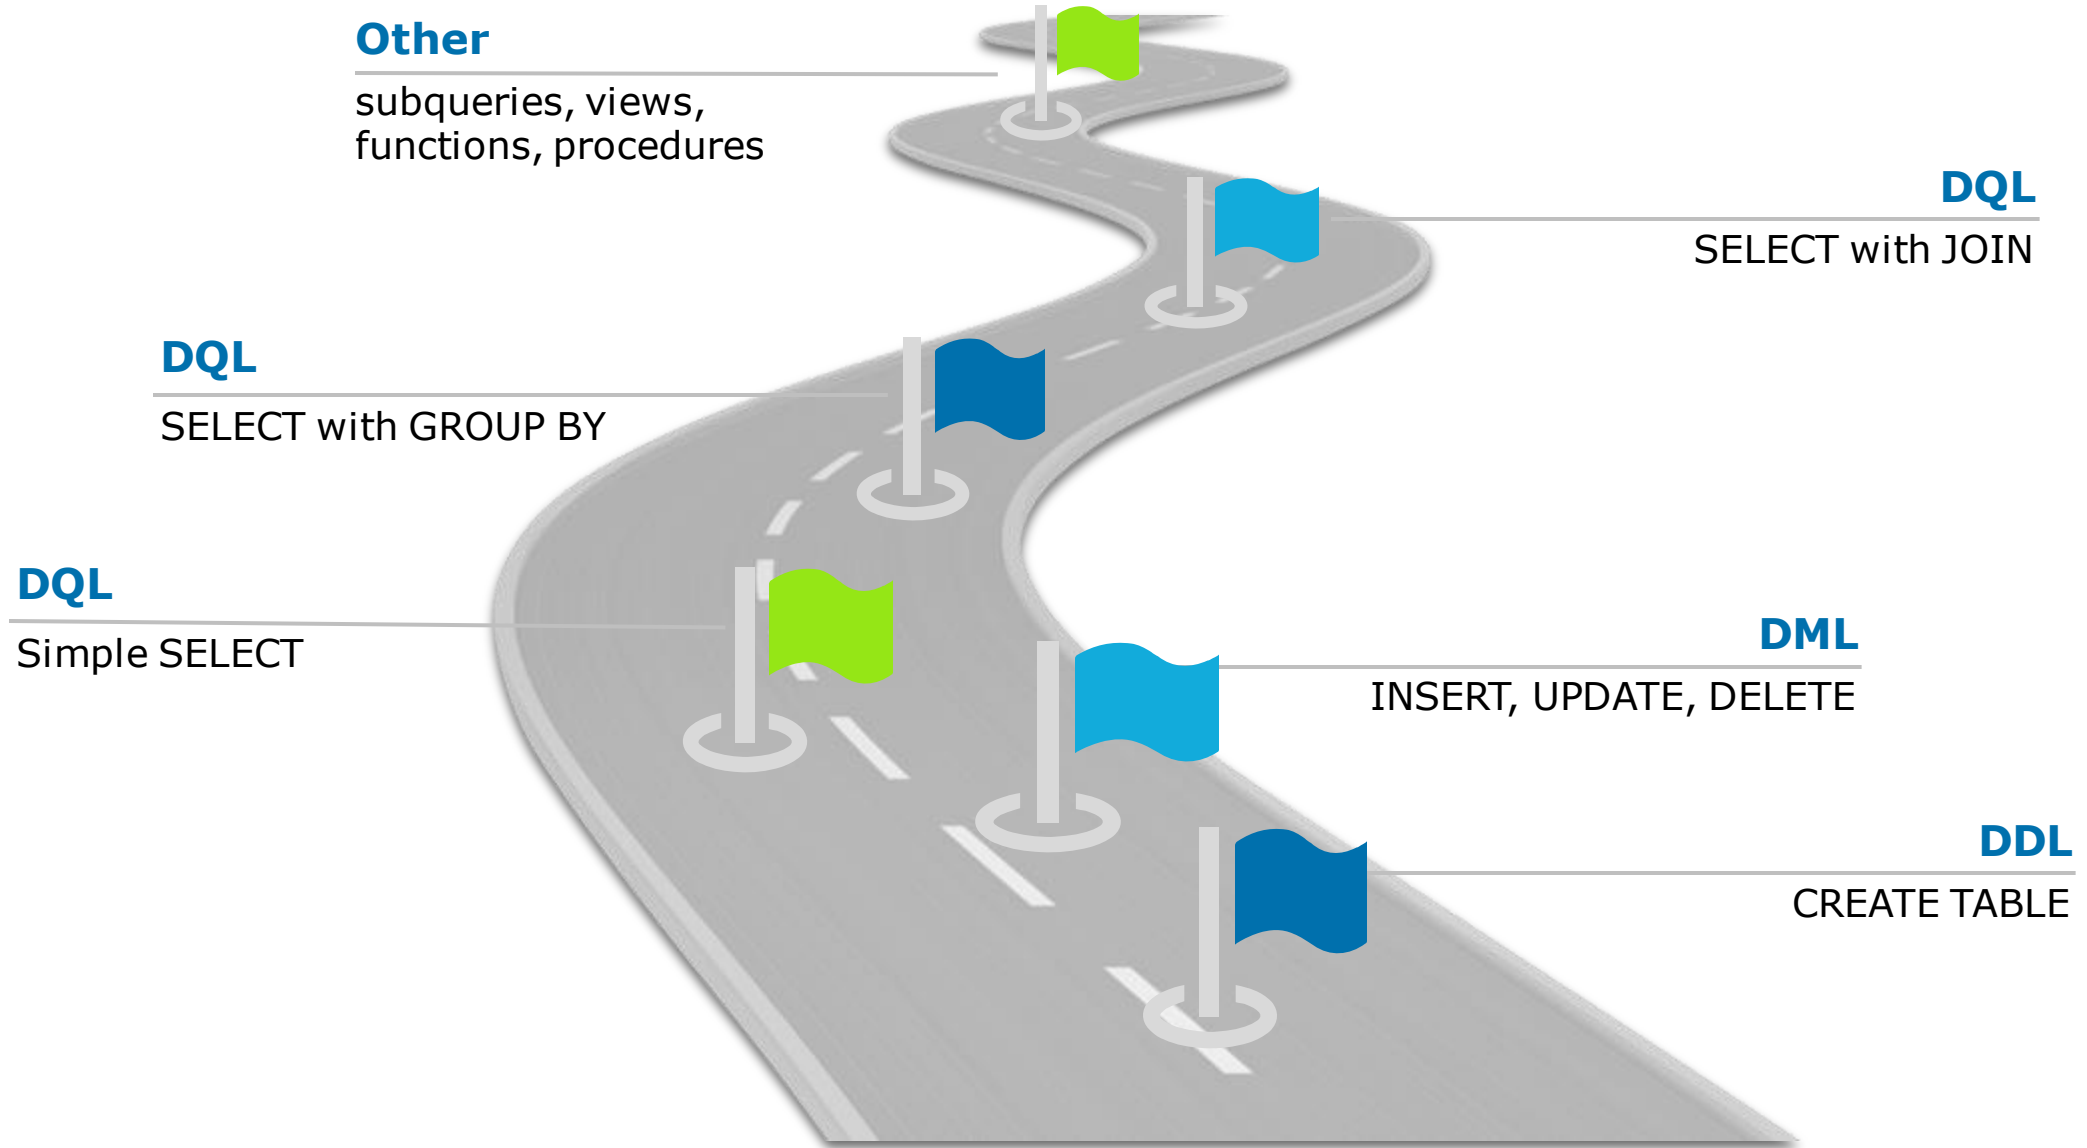
id	first_name	last_name	hired	created_by	created_at
1	Maciej	Kowalski	2017-01-01	hr1	2019-02-17
2	Armando	Roggio	2014-03-13	hr1	2019-01-01

id	first_name	last_name	hired	created_by	created_at
6	Steve	Jobs	1988-02-08	null	null

id	first_name	last_name	hired	created_by	created_at
3	Barack	Obama	2017-01-20	admin	2019-01-02
4	Bono	Vox	2014-07-13	admin	2019-01-01
5	Bill	Gates	2009-02-20	admin	2019-01-01
7	Larry	Ellison	2000-08-09	admin	2019-01-01


id	first_name	last_name	hired	created_by	created_at
8	Maciej	Nowak	2017-02-01	hr2	2017-02-01

Almost there...




Join - Types


SQL JOINS




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL  
OR B.Key IS NULL
```

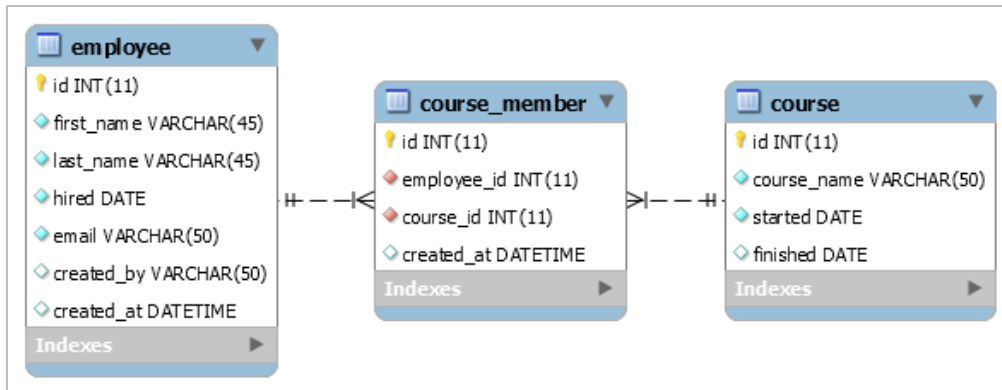
© C.L. Moffatt, 2008

<https://dev.mysql.com/doc/refman/8.0/en/join.html> ; Source: <https://pawelwiejkut.net/sql-on-conditions-in-left-join/>

Let's create more tables

```
CREATE TABLE jsk_db.course (  
    id INT NOT NULL AUTO_INCREMENT,  
    course_name VARCHAR(50) NOT NULL,  
    started DATE NOT NULL,  
    finished DATE DEFAULT NULL,  
    PRIMARY KEY (id),  
    UNIQUE INDEX uq_course_name (course_name)  
);
```

```
CREATE TABLE jsk_db.course_member (  
    id INT NOT NULL AUTO_INCREMENT,  
    employee_id INT NOT NULL,  
    course_id INT NOT NULL,  
    PRIMARY KEY (id),  
    UNIQUE INDEX uq_course_member  
        (employee_id, course_id),  
    CONSTRAINT fk_employee_id  
        FOREIGN KEY (employee_id)  
        REFERENCES employee (id),  
    CONSTRAINT fk_course_id  
        FOREIGN KEY (course_id)  
        REFERENCES course (id)  
);
```



JOIN

Table: employee

id	first_name	last_name	hired	created_by	created_at
1	Maciej	Kowalski	2017-01-01	hr1	2019-02-17
2	Armando	Roggio	2014-03-13	hr1	2019-01-01
3	Barack	Obama	2017-01-20	admin	2019-01-02
4	Bono	Vox	2014-07-13	admin	2019-01-01
5	Bill	Gates	2009-02-20	admin	2019-01-01
6	Steve	Jobs	1988-02-08	null	null
7	Larry	Ellison	2000-08-09	admin	2019-01-01
8	Maciej	Nowak	2017-02-01	hr2	2017-02-01

Table: course

id	course_name	started	finished
1	Deutschkurs 101	2019-02-20	null
2	Deutschkurs 102	2019-02-20	null
3	English course 101	2019-01-21	2019-01-31
4	English course 102	2019-02-20	null
5	English course 201	2019-02-20	

Table: course_member

id	employee_id	course_id	created_at
1	1	1	2019-02-20
2	1	2	2019-02-20
3	1	3	2019-02-20
4	1	4	2019-02-20
...

employee_id	first_name	last_name	Deutschkurs 101	Deutschkurs 102	English course 101	English course 102	English course 201
1	Maciej	Kowalski					
2	Armando	Roggio					
3	Barack	Obama					
4	Bono	Vox					
5	Bill	Gates					
6	Steve	Jobs					
7	Larry	Ellison					
8	Maciej	Nowak					

<https://dev.mysql.com/doc/refman/8.0/en/join.html>

Oracle syntax vs ANSI syntax

```
SELECT e.*  
FROM  
    employee e,  
    course_member cm  
WHERE  
    cm.employee_id = e.id;
```



```
SELECT e.*  
FROM employee e  
JOIN course_member cm ON (cm.employee_id = e.id);
```



DQL: SELECT - subqueries

```
SELECT c.*  
FROM course c  
WHERE c.id IN (  
    SELECT cm.course_id  
    FROM course_member cm  
    WHERE cm.created_at < '2019-02-01'  
);
```



```
SELECT c.*  
FROM course c  
WHERE EXISTS (  
    SELECT 1  
    FROM course_member cm  
    WHERE cm.course_id = c.id  
    AND cm.created_at < '2019-02-01'  
);
```



```
SELECT c.*  
FROM comments c  
JOIN properties p ON (p.comments_id = c.id)  
WHERE p.property = 'OLD';
```



<https://dev.mysql.com/doc/refman/8.0/en/subqueries.html> ; <https://dev.mysql.com/doc/refman/8.0/en/optimizing-subqueries.html> ; <https://dev.mysql.com/doc/refman/8.0/en/subquery-optimization-with-exists.html>

VIEW

```
CREATE OR REPLACE VIEW v_employees AS  
SELECT e.id, e.first_name, e.last_name  
FROM employee e;
```

```
SELECT e.* FROM v_employees e;
```

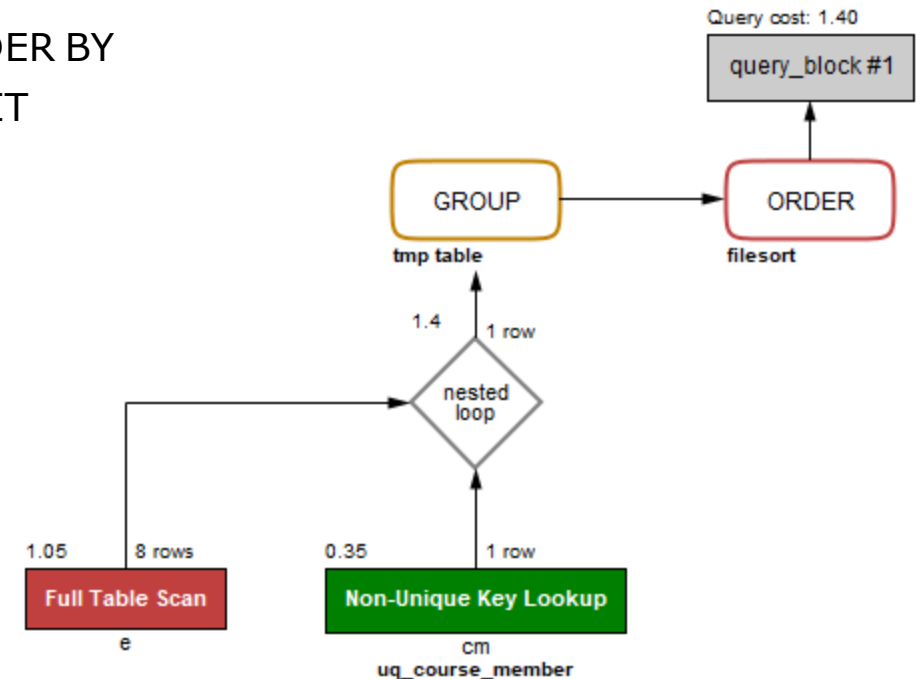
```
CREATE OR REPLACE VIEW v_courses_and_attendants  
SELECT c.id AS course_id,  
       c.course_name,  
       COUNT(cm.id) AS num_participants  
FROM course c  
LEFT JOIN course_member cm ON (cm.course_id = c.id)  
GROUP BY c.id, c.course_name  
ORDER BY c.course_name ASC;
```

```
SELECT c.*,  
       v.num_participants  
FROM course c  
JOIN v_courses_and_attendants v  
  ON (v.course_id = c.id)  
WHERE v.num_participants > 2;
```

Order of execution

```
SELECT e.id, e.first_name, e.last_name,  
       COUNT(cm.id) AS num_courses  
FROM employee e  
JOIN course_member cm ON (cm.employee_id = e.id)  
WHERE e.first_name LIKE '%a%'  
GROUP BY e.id, e.first_name, e.last_name  
HAVING COUNT(*) > 1  
ORDER BY e.first_name ASC, e.last_name ASC;
```

1. FROM & JOIN
2. WHERE
3. GROUP BY & aggregate functions
4. HAVING
5. WINDOW functions
6. ORDER BY
7. LIMIT



Functions and procedures

- Functions – processing input data and returns result / modify function's parameters
- Procedures – processing data / returns rowset

```
DROP FUNCTION IF EXISTS fn_hello;

DELIMITER $$

CREATE FUNCTION fn_hello(name VARCHAR(50))
    RETURNS VARCHAR(100) DETERMINISTIC
BEGIN
    RETURN CONCAT('Hello ', name, '!');
END
$$

DELIMITER ;
```

```
-- Sample call
SELECT fn_hello('Jan');

-- Use with data from table
SELECT e.id,
       e.first_name,
       e.last_name,
       fn_hello(e.first_name) AS hello
FROM employee e
ORDER BY e.first_name ASC;
```

<https://dev.mysql.com/doc/refman/8.0/en/create-procedure.html>

Why it is worth to alias tables/columns?

```
SELECT DISTINCT
    serial_no AS BARCODE_NO
  , org_id AS ORG_ID
  , wip_entity_id AS WOID
  , model_suffix AS MTRLID
  , model AS MODLID
  , suffix AS SFFX_NAME
  , line_id AS PCSGID
  , buyer_code AS CUSTOMERID
  , serial AS LBL_SERIAL_NO
  , quantity
  , wip_entity_name
  , serial_no AS BUYER_SERIAL_NO
  , TO_CHAR(sysdate, 'YYYYMMDD') AS PUB_YMD
  , TO_TIMESTAMP(SYSDATE, 'YYYY-MM-DD HH24:MI:SS') AS CREATED_AT
FROM prodplan.daily_prod_plan pln
JOIN matmgmt.workorder wo ON (WOID = wip_entity_id)
LEFT JOIN matmgmt.serialnumbers snb ON (
    division_code = division
  AND work_order = mfg_order)
LEFT JOIN matmgmt.serialnumbers_sent sns ON (
    division_code = division
  AND work_order = mfg_order
  AND serial = serial)
LEFT JOIN prodplan.xxx_register_inf lph ON (barcode_no = serial_no)
WHERE 1=1

    AND org_id = '1111'
    AND division = 'ABC'
    AND item_type = 'F'
    AND barcode_no is null
    AND mfg_order = '111X0123'
    AND serial = 'KP001'
    AND ROWNUM = 1;
```



```
SELECT DISTINCT
    snb.serial_no AS BARCODE_NO
  , pln.org_id AS ORG_ID
  , wo.wip_entity_id AS WOID
  , wo.model_suffix AS MTRLID
  , wo.model AS MODLID
  , wo.suffix AS SFFX_NAME
  , pln.line_id AS PCSGID
  , wo.buyer_code AS CUSTOMERID
  , snb.serial AS LBL_SERIAL_NO
  , wo.quantity
  , wo.wip_entity_name
  , sns.serial_no AS BUYER_SERIAL_NO
  , TO_CHAR(sysdate, 'YYYYMMDD') AS PUB_YMD
  , TO_TIMESTAMP(SYSDATE, 'YYYY-MM-DD HH24:MI:SS') AS CREATED_AT
FROM prodplan.daily_prod_plan pln
JOIN matmgmt.workorder wo ON (pln.WOID = wo.wip_entity_id)
LEFT JOIN matmgmt.serialnumbers snb ON (
    snb.division_code = wo.division
  AND snb.work_order = pln.mfg_order)
LEFT JOIN matmgmt.serialnumbers_sent sns ON (
    sns.division_code = wo.division
  AND sns.work_order = pln.mfg_order
  AND sns.serial = snb.serial)
LEFT JOIN prodplan.xxx_register_inf lph ON (lph.barcode_no = snb.serial_no)
WHERE 1=1

    AND pln.org_id = '1111'
    AND wo.division = 'ABC'
    AND wo.item_type = 'F'
    AND lph.barcode_no is null
    AND pln.mfg_order = '111X0123'
    AND snb.serial = 'KP001'
    AND ROWNUM = 1;
```



Window Functions (MySQL) / Analytic Functions (Oracle)

```

1  mysql> SELECT
2      val,
3      ROW_NUMBER() OVER w AS 'row_number',
4      RANK() OVER w AS 'rank',
5      DENSE_RANK() OVER w AS 'dense_rank'
6  FROM numbers
7  WINDOW w AS (ORDER BY val);
8
9  +-----+-----+-----+-----+
10 | val | row_number | rank | dense_rank |
11 +-----+-----+-----+-----+
12 | 1 | 1 | 1 | 1 |
13 | 1 | 2 | 1 | 1 |
14 | 2 | 3 | 3 | 2 |
15 | 3 | 4 | 4 | 3 |
16 | 3 | 5 | 4 | 3 |
17 | 3 | 6 | 4 | 3 |
18 | 4 | 7 | 7 | 4 |
19 | 4 | 8 | 7 | 4 |
20 | 5 | 9 | 9 | 5 |
21 +-----+-----+-----+-----+

```

MySQL: <https://dev.mysql.com/doc/refman/8.0/en/window-function-descriptions.html> ; Oracle: https://docs.oracle.com/cd/E11882_01/server.112/e41084/functions004.htm#SQLRF06174

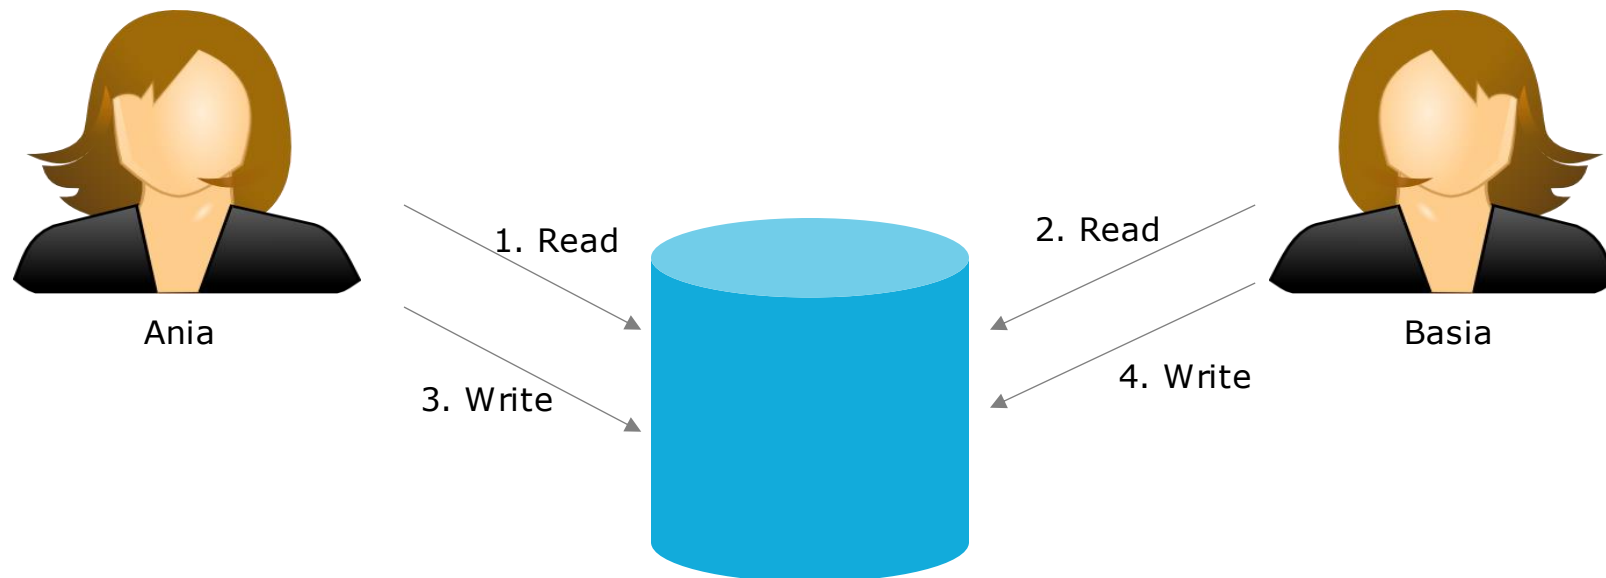
Agenda



- Before we Begin
- Relational Databases (DB)
- Structured Query Language (SQL)
- **Data locking strategies**
- Advanced Databases Topics
- Modeling databases
- Assignments


No locking

- Effect: Last one wins.
- Problematic from the business point of view.



Pesimistic locking

- Effect: First one locks the data.
- Can result in long lasting locks that disturb the work.




Optimistic locking

- Effect: Only the first write is successful.
- Standard approach

Implementation:

1. Table has an attribute: version (int)
2. When you update the row, you increment the version
3. You can only update if row's version matches Yours's version

```
UPDATE table  
SET attribute = new_value  
  , version = version+1  
WHERE id = object_id  
  AND version = old_version;
```




Agenda



- Before we Begin
- Relational Databases (DB)
- Structured Query Language (SQL)
- Data locking strategies
- **Advanced Databases Topics**
- Modeling databases
- Assignments

Transactions. ACID Properties



Isolation levels:

1. Serializable
2. Repeatable read
3. **Read committed**
4. Read uncommitted

Transactions

Connection 1

SHOW VARIABLES LIKE 'transaction_isolation';
-- REPEATABLE-READ

BEGIN; -- start transaction
INSERT INTO test (comment) VALUES ('test1');

INSERT INTO test (comment) VALUES ('test2');
COMMIT; -- or ROLLBACK

SELECT * FROM test;
-- test1 not visible

SELECT * FROM test;
-- test1 and test2 visible

Connection 2

Agenda



- Before we Begin
- Relational Databases (DB)
- Structured Query Language (SQL)
- Data locking strategies
- Advanced Databases Topics
- **Modeling databases**
- Assignments

Consider this simple exercise




Description of the domain:

1. The database is to be used for personal movie organization purposes.
2. All movies are stored on some kind of a data carrier (e.g. VideoCD, DVD, BlueRay).
3. More than one movie can be stored on a carrier.
4. Each carrier is identified by a unique number.
5. The database must contain a photo of the carriers front cover.
6. For each and every movie a title, year of production and one director must be present and all actors must be listed.
7. ...

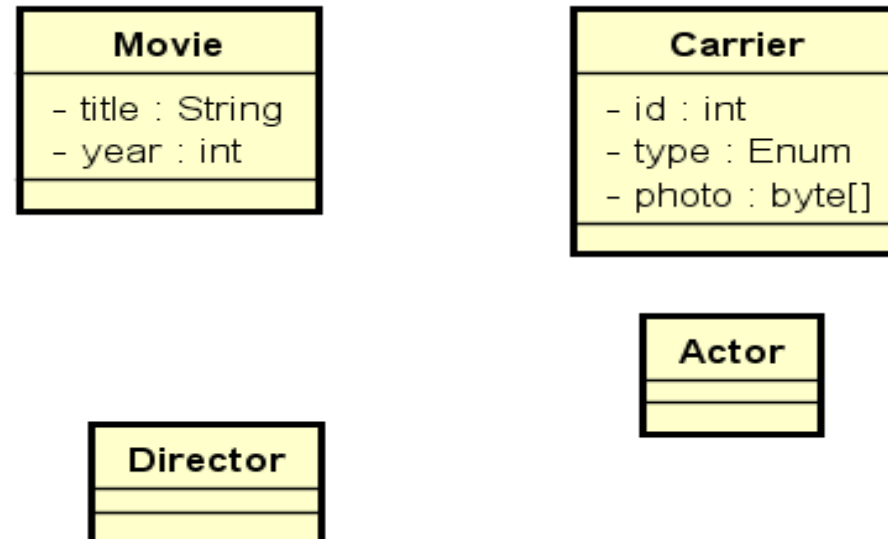
Identify all entities

The database is to be used for personal movie organization purposes. All **movies** are stored on some kind of a data **carrier** (e.g. VideoCD, DVD, BlueRay), more than one **movie** can be stored on a **carrier**. Each **carrier** is identified by a unique number. The database must contain a photo of the **carriers** front cover. For each and every movie a title, year of production and one **director** must be present and all **actors** must be listed. ...



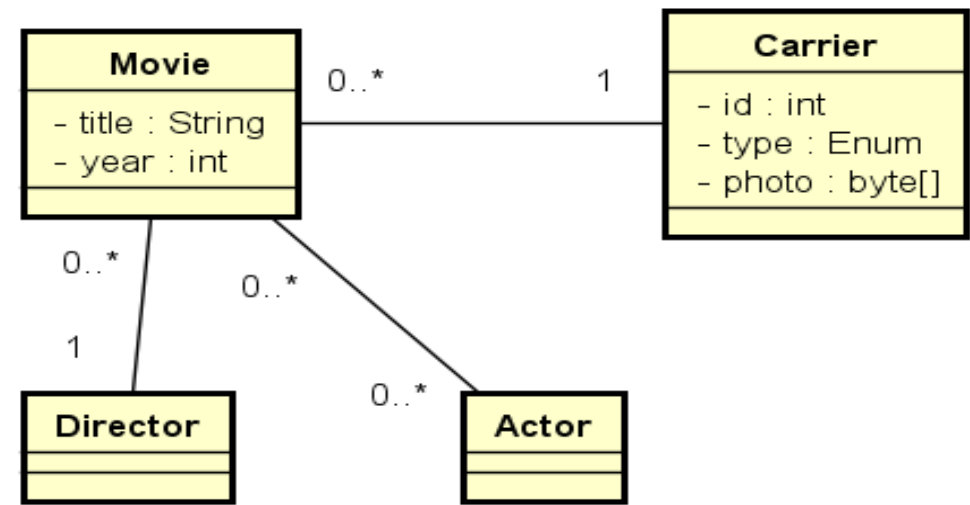
Identify entity attributes

The database is to be used for personal movie organization purposes. All **movies** are stored on some kind of a data **carrier** (e.g. VideoCD, DVD, BlueRay), more than one **movie** can be stored on a **carrier**. Each **carrier** is identified by a **unique number**. The database must contain a **photo of the carriers front cover**. For each and every movie a **title**, **year of production** and one **director** must be present and all **actors** must be listed. ...

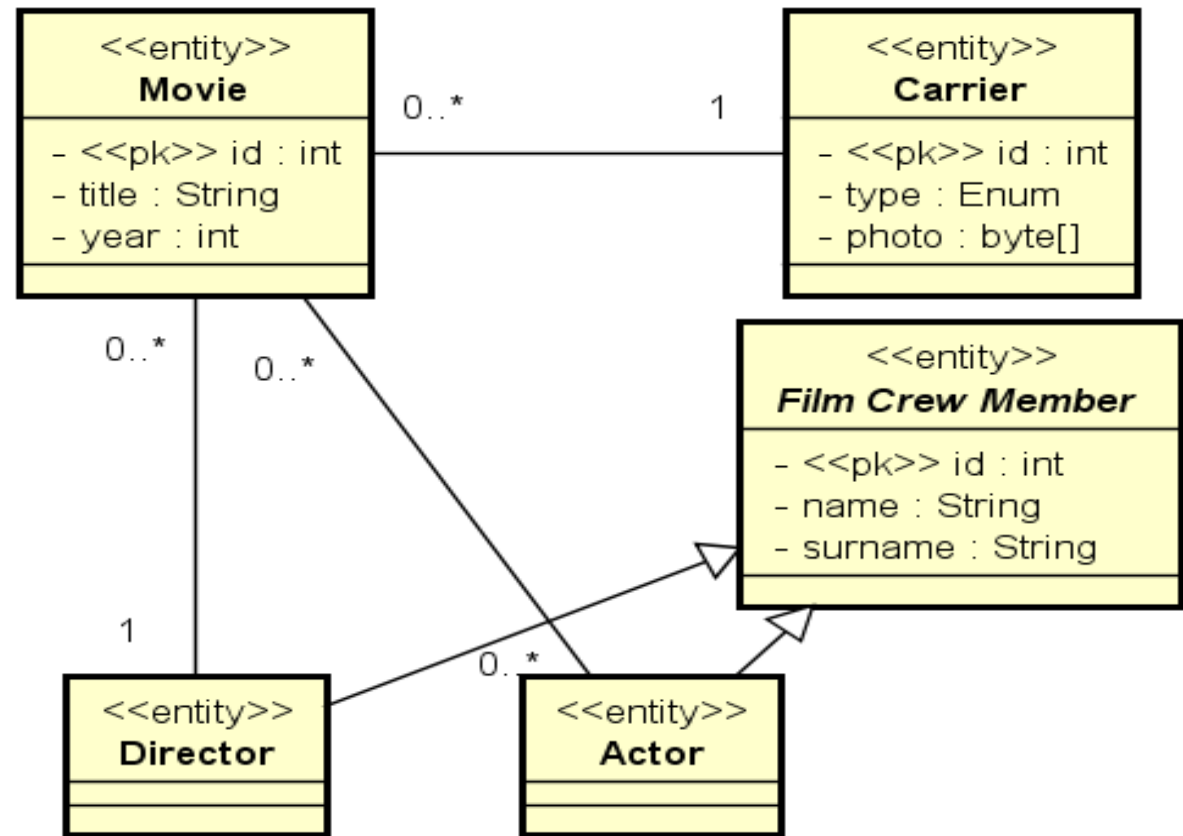


Identify the relations between the entities

The database is to be used for personal movie organization purposes. **All movies** are stored on some kind of a data **carrier** (e.g. VideoCD, DVD, BlueRay), **more than one movie** can be stored on a **carrier**. Each **carrier** is identified by a **unique number**. The database must contain a **photo of the carriers front cover**. For each and every movie a **title**, **year of production** and **one director** must be present and **all actors** must be listed.



Validate, normalize, do a sanity check.
 What you get is a business database model.



Create a technical database model

<<table>> Movie	
- <<pk>> id : int	
- title : String	
- year : int	
- carrier_id : int	
- director_id : int	

<<table>> Carrier	
- <<pk>> id : int	
- type : Enum	
- photo : byte[]	

Movie2Actor	
- <<pk>> movie_id : int	
- <<pk>> actor_id : int	

<<table>> Film Crew Member	
- <<pk>> id : int	
- name : String	
- surname : String	

Agenda



- Before we Begin
- Relational Databases (DB)
- Structured Query Language (SQL)
- Data locking strategies
- Advanced Databases Topics
- Modeling databases
- **Assignments**

Assignments



1. Proszę zwrócić uwagę na informacje zawarte w mailu powitalnym
 1. zwróćcie uwagę na sugestie
 2. pomocne linki
 3. pytajcie trenerów!
2. Niektóre zadania pozwalają na interpretacje
 1. proponujcie rozwiązania, argumentujcie „dlaczego tak a nie inaczej”
 2. doprecyzujcie niektóre tematy na konsultacjach
 3. brońcie Waszego rozwiązania podczas konsultacji
3. Przydatny plik do analizy: „JSK DB, live sample*.sql”

Things to remember



1. Carefully read questions
2. Listen Your trainer
3. Assignments:
 1. Use aliases on tables / subqueries
 2. Do not use subqueries in places, where they are inefficient
 3. Sample scripts were prepared in a way, in which they could be easily re-created; do the same



SEC1



Questions?



People matter, results count.

This presentation contains information that may be privileged or confidential and is the property of the Capgemini Group.

Copyright © 2020 Capgemini. All rights reserved.

About Capgemini

Capgemini is a global leader in consulting, digital transformation, technology and engineering services. The Group is at the forefront of innovation to address the entire breadth of clients' opportunities in the evolving world of cloud, digital and platforms. Building on its strong 50-year+ heritage and deep industry-specific expertise, Capgemini enables organizations to realize their business ambitions through an array of services from strategy to operations. Capgemini is driven by the conviction that the business value of technology comes from and through people. Today, it is a multicultural company of 270,000 team members in almost 50 countries. With Altran, the Group reported 2019 combined revenues of €17billion.

Learn more about us at

www.capgemini.com